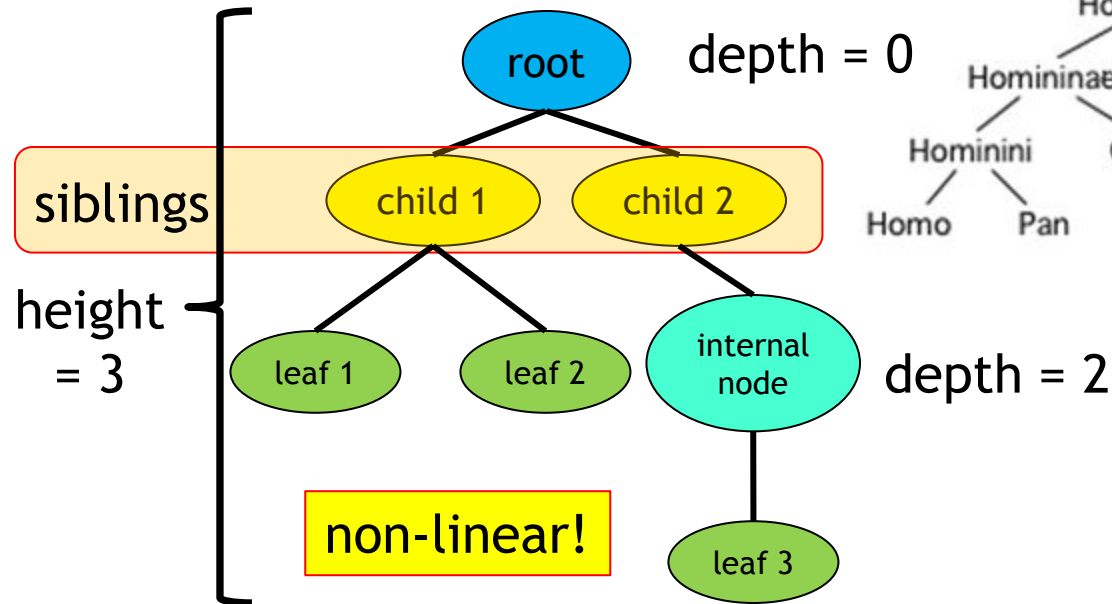



Trees


- A data structure for *hierarchical data*
 - data stored in **nodes** (or **vertices**)



Modern Hominoid Classification



- Every node except the root has exactly one parent



Vocabulary

1. Root node
2. Leaf node
3. Edge
4. Child/parent
5. Siblings
6. Internal node
7. Ancestor/descendent
8. Depth of a node
9. Height of a tree


Binary Trees

- A **binary tree** is a tree where each internal node has at most two children
 - sometimes indicated as left-child and right-child
- Implementation:

```
typedef struct node {  
    int data; /* or some other type */  
    struct node *left;  
    struct node *right;  
} tNode;
```
- If an internal node only has one child, `left` or `right` will be `NULL`
 - for a leaf, `left = right = NULL`
- To access: `tNode *root; /* NULL, if tree is empty */`

Creating a node

- Suppose we want to create a new node containing some data D
 - the left child should be the sub-tree L
 - the right child should be the sub-tree R
 - L and R could be empty




```
tNode *makeNode(int D, tNode *L, tNode *R) {  
    tNode *node = (tNode *) malloc(sizeof(tNode));  
    if(node == NULL) { return NULL; }  
    node->data = D;  
    node->left = L;  
    node->right = R;  
    return node;  
}
```

Question

- What kind of a tree does the following code make?

```
tNode *makeNode(int D, tNode *L, tNode *R) {  
    /* code from previous slide */  
}
```



```
void main() {  
    tNode *n1, *n2, *n3, *n4, *root;  
    n1 = makeNode(1, NULL, NULL);  
    n2 = makeNode(2, n1, NULL);  
    n3 = makeNode(3, NULL, NULL);  
    n4 = makeNode(4, n3, n2);  
    root = makeNode(0, NULL, n4);  
}
```



Simplified forall function for trees

- Suppose we want to print all nodes in a tree:

```
void forall(tNode *node) {  
    if(node == NULL) { return; }  
    printf("%d\n", node->data);  
    forall(node->left);  
    forall(node->right);  
}
```





- Question 1: What happens on this?
- Question 2: How about with this change?
 - first version: **preorder** traversal (this, left, right)
 - second version: **inorder** traversal (left, this, right)
 - also: **postorder** traversal (left, right, this)

Simplified forall function for trees

- Suppose we want to print all nodes in a tree:

```
void forall(tNode *node) {  
    if(node == NULL) { return; }  
    printf("%d\n", node->data);  
    forall(node->left);  
    forall(node->right);  
}
```



- Question 1: What happens on this?
- Question 2: How about with this change?
 - first version: **preorder** traversal (this, left, right)
 - second version: **inorder** traversal (left, this, right)
 - also: **postorder** traversal (left, right, this)


General preorder

```
tNode *preorder(tNode *node,
                int(*task)(int, void*), void*),
                void *args) {
    if(node == NULL) { return NULL; }
    if(!task(node->data, args)) { return node; }
    tNode *temp = preorder(node->left, task, args);
    if(temp != NULL) { return temp; }
    return preorder(node->right, task, args);
}
```

Question

- What does this function compute?

```
int strange(tNode *node) {  
    if(node == NULL) { return 0; }  
    return strange(node->left)  
        + node->data  
        + strange(node->right);  
}
```




- For the example:* $0 + 3 + 4 + 1 + 2$
 - in general, sum of data
- Question:** The above function uses inorder traversal. What would happen if we used preorder traversal? Postorder?

Freeing memory

- The following function frees memory used by a tree:

```
void freeTree(tNode *node) {  
    if(node == NULL) { return; }  
    freeTree(node->left);  
    free(node);  
    freeTree(node->right);  
}
```




- **Question:** The above function uses postorder traversal. What would happen if we used inorder traversal? Preorder?

An application: expression trees


- Evaluating the tree:

```
// assume node != NULL
int eval(tNode *node) {
    if(isInt(node->data)) {
        return toInt(node->data);
    }
    if(isPlus(node->data)) {
        return eval(tNode->left)
            + eval(tNode->right);
    }
    if(isMult(node->data)) { ... }
    if(isUMinus(node->data)) {
        return -eval(node->left);
    }
    // never reach here
}
```



Ordered Binary Trees


- Suppose the data-type used in tree nodes supports ordering
 - e.g., any two `ints` can be compared using `<`
- We can then define an **ordered binary tree** where every node x satisfies this ordering property:
 - the data at all nodes in x 's left sub-tree is less than or equal to the data at node x
 - the data at all nodes in x 's right sub-tree is greater than the data at node x
- Finding data in a tree
 - without the ordering property
 - with the ordering property



Finding data in an ordered tree

- This function assumes that `node` is the root of an ordered tree


```
int find(tNode *node, int D) {  
    if(node == NULL) { return 0; }  
    if(D == node->data) {  
        return 1;  
    } else if(D < node->data) {  
        return find(node->left, D);  
    } else {  
        return find(node->right, D);  
    }  
}
```



Finding data in an ordered tree

- This function assumes that `node` is the root of an ordered tree


```
int find(tNode *node, int D) {  
    if(node == NULL) { return 0; }  
    if(D == node->data) {  
        return 1;  
    } else if(D < node->data) {  
        return find(node->left, D);  
    } else {  
        return find(node->right, D);  
    }  
}
```



Inserting data into an ordered tree

- Our ordered tree data structure will not permit duplicate data
 - can be modified to allow duplicates

```
tNode *insert(tNode *node, int D) {  
    if(node == NULL) {  
        return makeNode(D, NULL, NULL);  
    }  
    if(D < node->data) {  
        node->left =  
            insert(node->left, D);  
    } else if(D > node->data) {  
        node->right =  
            insert(node->right, D);  
    }  
    return node;  
}
```




Understanding insert

- Draw the tree that will be created

```
tNode *insert(tNode *node, int D) {  
    /* code from previous slide */  
}
```

```
void main() {  
    tNode* root =  
        insert(  
            insert(  
                insert(NULL, 2),  
                    3),  
                1);  
}
```



Question 1: Worst-case running time?
Question 2: What if `makeNode` fails?

A better insert

```
tNode *insert(tNode *node, int D) {  
    if(find(node, D)) { return NULL; }  
    tNode *newNode = makeNode(D, NULL, NULL);  
    if(newNode == NULL) { return NULL; }  
    return rec_ins(node, newNode);  
}
```

```
tNode *rec_ins(tNode *node, tNode *newNode) {  
    if(node == NULL) { return newNode; }  
    if(newNode->data < node->data) {  
        node->left = rec_ins(node->left, newNode);  
    } else if(newNode->data > node->data) {  
        node->right = rec_ins(node->right, newNode);  
    }  
    return node;    If newNode is the root of an ordered tree,  
                    will rec_ins merge the two ordered trees?  
}
```


Merging two trees

- Write a function to merge all nodes from `src` into `dest`

```
tNode *merge(tNode *src, tNode *dest) {  
    if(src == NULL) { return dest; }  
    dest = insert(dest, src->data);  
    dest = merge(src->left, dest);  
    return merge(src->right, dest);  
}
```

- The above code traverses `src` nodes in preorder
 - what if we used inorder? Postorder?
- If both trees have n nodes initially, how long will `merge` take?

Deleting data from an ordered tree

- Recall that our ordered trees do not permit duplicates

```
tNode *delete(tNode *node, int D) {  
    if(node == NULL) { return NULL; }  
    if(D < node->data) {  
        node->left = delete(node->left, D);  
    } else if(D > node->data) {  
        node->right = delete(node->right, D);  
    } else {  
        tNode *oldNode = node;  
        node = merge(node->left, node->right);  
        freeTree(oldNode->left);  
        free(oldNode);  
    }  
    return node;  
}
```

Slow because it:

- does not recycle old nodes
- wastefully compares data in `node->left`, `node->right` subtrees

What does this code do?

```
tNode *extractMax(tNode *node, tNode *ans) {
    // assume node != NULL
    if(node->right == NULL) {
        ans->data = node->data;
        tNode *temp = node->left;
        free(node);
        return temp;
    }
    node->right = extractMax(node->right, ans);
    return node;
}
```

```
tNode *extractMin(tNode *node, tNode *ans) {
    /* similar */
}
```

A better delete

```
tNode *delete(tNode *node, int D) {
    if(node == NULL) { return NULL; }
    if(D < node->data) {
        node->left = delete(node->left, D);
    } else if(D > node->data) {
        node->right = delete(node->right, D);
    } else {
        if(node->left != NULL)
            node->left = extractMax(node->left, node);
        else if(node->right != NULL)
            node->right = extractMin(node->right, node);
        else { free(node); return NULL; }
    }
    return node;
}
```

Exam 1: Review

- Question 1 (4 points): What is the error?

```
1: void freeList(List list) {
2:     Node *this = (Node *) list;
3:     if(this != NULL)
4:         freeList(this->next);
5:     free(this);
6: }
```

Bad typecast?
No base-case?
Type mismatch?
Segmentation fault?

- Question 2 (6 points): Worst-case time complexity for size n list:

ListEntry previous(List list, ListEntry e);

ArrayList	$O(1)$
LinkedList	$O(\log n)$
	$O(n)$

Question 3 (10 points)

```
int sizeof(DLL_node *node) {  
    // assume node != NULL;  
    /* Idea: start from node and go  
       backwards until NULL, and go  
       forwards until NULL, counting  
       as you go  
    */  
    int ans = 1;  
    for(DLL_node *p=node->prev; p != NULL; p = p->prev)  
        ans++;  
    for(DLL_node *p=node->next; p != NULL; p = p->next)  
        ans++;  
    return ans;  
}
```

Question 4 (5+5 points)

```
int peek(Queue); // returns data at front of queue
```

```
first = peek(myQueue);
```

- No way to signal failure on empty queue

```
int peek(Queue q, int *peeked_value);
```

```
if(peek(myQueue, &first)) { ... }
```


```
int peek(Queue q, int *error_flag);
```

```
first = peek(myQueue, &ok);
```

```
if(ok) { ... }
```

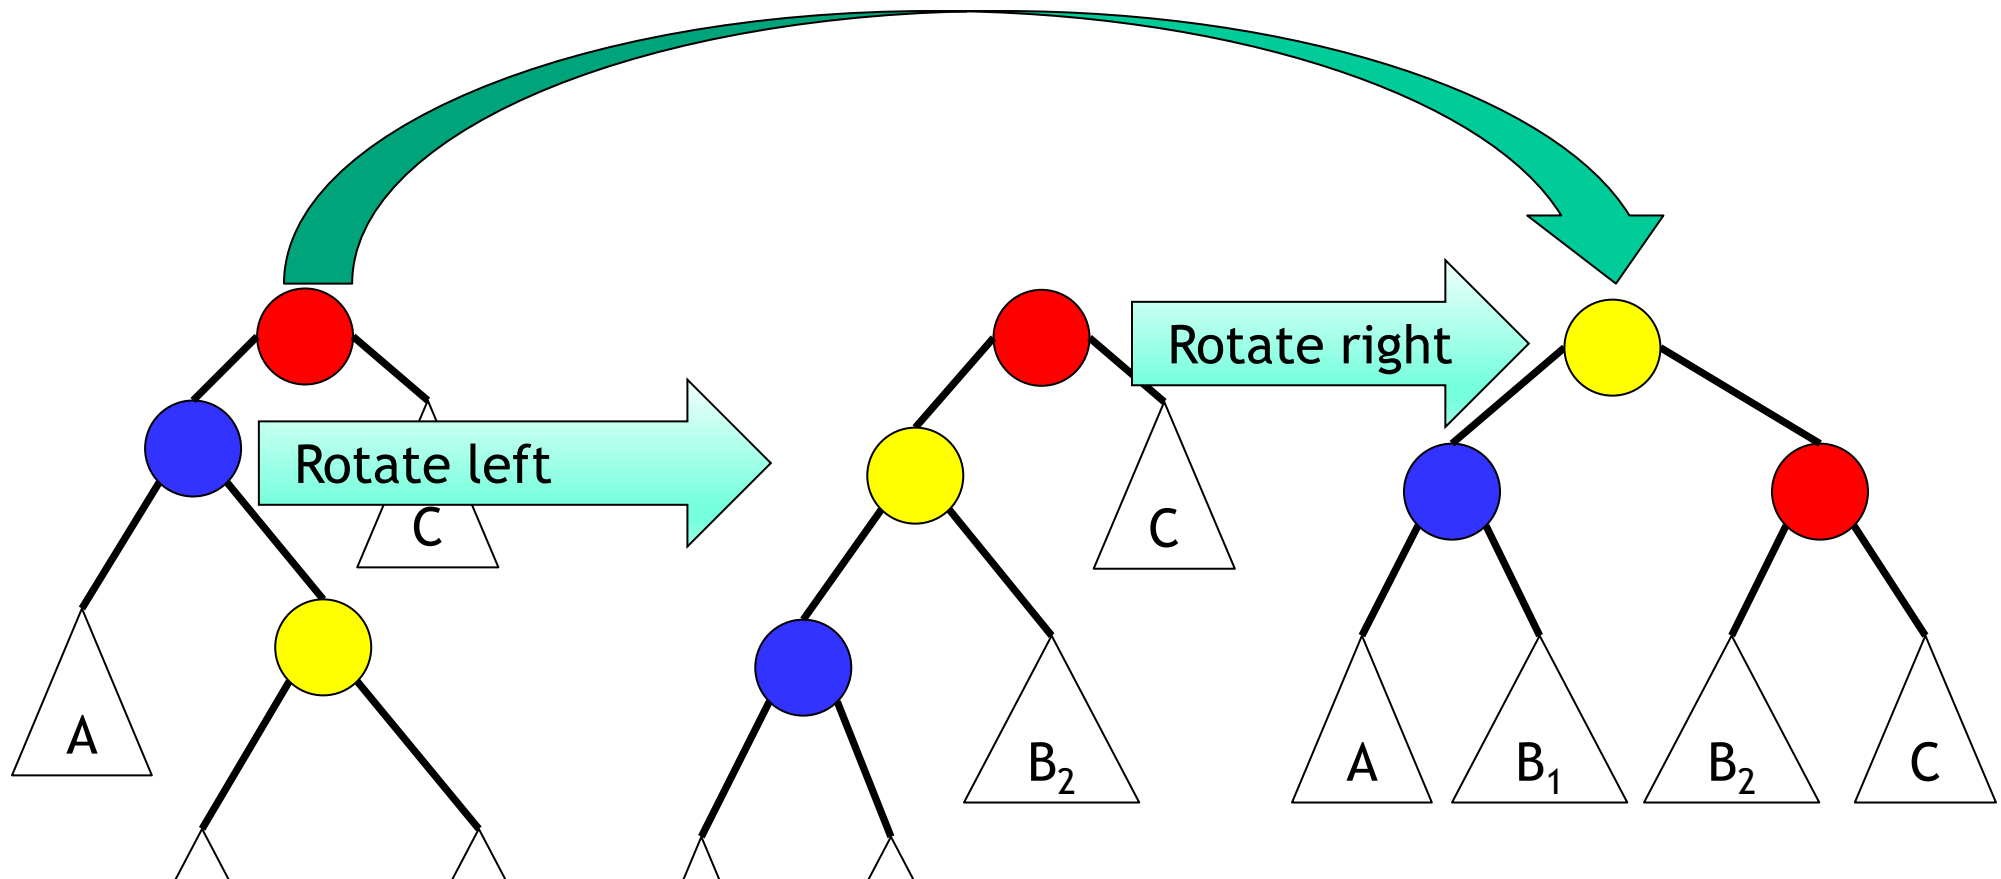
Balanced Search Trees

- How long does it take to determine whether a particular value x appears in an ordered tree T containing n nodes?
 - Worst case time is $O(h)$, where h is the height of T
 - Minimum height (T is balanced): $h = \lfloor \log_2 n \rfloor$
 - Maximum height (T is lopsided): $h = n - 1$
- Key idea: Keep the tree balanced (prevent it from getting lopsided)



Dealing with Left Right and Right Left

- Do a double-rotation!



[PDF] SOVIET MATHEMATICS DOKLADY

professor.ufabc.edu.br/~jesus.mena/courses/mc3305-2q.../AED2-10-avl-paper.pdf

by GM Adel'son-Vel'skii 1962 - Cited by 4 Related articles

An algorithm for the organization of information by G. M. Adel'son-Vel'skii and E. M. Landis. Soviet Mathematics Doklady, 3, 1259-1263, 1962.

sure that height
between left and
s is at most 1

AVL Tree Implementation

```
typedef struct node {  
    int data; /* or some other type */  
    struct node *left;  
    struct node *right;  
    int height;  
} tNode;
```

```
int h(tNode *node) {  
    return node == NULL ? -1 : node->height;  
}
```

```
tNode *makeNode(int D, tNode *L, tNode *R) {  
    tNode *node = (tNode *) malloc(sizeof(tNode));  
    ...  
    node->height = 1 + MAX(h(L), h(R));  
    return node;  
}
```

Insert in an AVL tree

```
tNode *insert(tNode*node,int D) {newNode) {
    if (find(node,D)) { return NULL; }
    tNode *newNode = makeNode(D, NULL, NULL);
    if (newNode->data < node->data) {
        if (newNode == NULL) { return NULL; }
        node->left = rec_ins(node->left, newNode);
        return rec_ins(node, newNode);
    }
    if (h(node->left) > h(node->right)) {
        node->height = h(node->left) + 1;
    }
} else if (newNode->data > node->data) {
    node->right = rec_ins(node->right, newNode);
    if (h(node->right) > h(node->left)) {
        node->height = h(node->right) + 1;
    }
}
return balance(node);
}
```

Balancing nodes in an AVL tree

```
tNode *balance(tNode *node) {
    if(node == NULL) { return NULL; }
    if(h(node->left) > h(node->right)+1) {
        if(h(node->left->left) > h(node->left->right) {
            return handleLL(node);
        } else {
            return handleLR(node);
        }
    } else if(h(node->right) > h(node->left)+1) {
        if(h(node->right->right) > h(node->right->left)) {
            return handleRR(node);
        } else {
            return handleRL(node);
        }
    } else { return node; }
}
```

handleLL


```
tNode *handleLL(tNode *node) {  
    tNode *A, *B, *C, *result;  
    result = node->left;  
    A = result->left;  
    B = result->right;  
    C = node->right;  
    result->right = node;  
    node->left = B;  
    node->height = 1 + MAX(h(B), h(C));  
    result->height = 1 + MAX(h(A), node->height);  
    return result;  
}
```

$O(1)$ time

Hence, `balance()` is also $O(1)$ time

Thus, insert is $O(h)$

We will prove that h is $O(\log n)$



Delete in an AVL tree

```
tNode *delete(tNode *node, int D) {
    if(node == NULL) { return NULL; }
    if(D < node->data)
        node->left = delete(node->left, D);
    else if(D > node->data)
        node->right = delete(node->right, D);
    else {
        if(node->left != NULL)
            node->left = extractMax(node->left, node);
        else if(node->right != NULL)
            node->right = extractMin(node->right, node);
        else { free(node); return NULL; }
    }
    node->height = 1 + MAX(h(node->left), h(node->right));
    return balance(node);
}
```

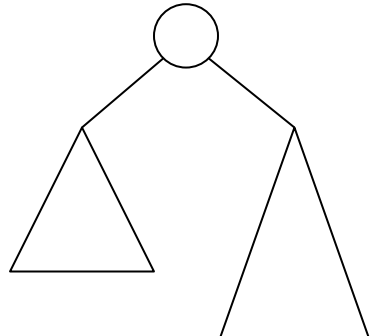
Space overhead per node

- Our AVL tree nodes store the (absolute) height of each node in the tree

```
typedef struct node {  
    int data;  
    struct node *left;  
    struct node *right;  
    int height; // sizeof(int) overhead per node  
} tNode;
```

- This can be reduced to two bits overhead per node:
 - LeftHeavy, Balanced, RightHeavy
- Slightly more complex insert/delete/balance
 - another example of the space/time tradeoff!

AVL trees with n nodes have $O(\log n)$ height

- **Recall:** AVL Trees ensure that height difference between left and right sub-trees is at most 1
- Let $m(h)$ be the minimum number of nodes in an AVL tree of height h
 - If an AVL tree with n nodes has height h then $n \geq m(h)$
 - Hence, $h \leq m^{-1}(n)$ [assuming m is an increasing function]
- Now $m(0) = 1$, $m(1) = 2$, $m(2) = 4$, $m(3) = 1 + m(2) + m(1) = 7$ 
- $\forall h \geq 2$, $m(h) = 1 + m(h-1) + m(h-2)$
 $F(0) = F(1) = 1$; $\forall h \geq 2$, $F(h) = F(h-1) + F(h-2)$
 - Prove that $\forall h \geq 0$, $m(h) = F(h+2) - 1$

Critique this “proof”

- **Given:** $m(0) = 1, m(1) = 2; \forall h \geq 2, m(h) = 1 + m(h-1) + m(h-2)$
and $F(0) = F(1) = 1; \forall h \geq 2, F(h) = F(h-1) + F(h-2)$
- We will prove that $\forall h \geq 0, m(h) = F(h+2) - 1$ by induction on h
- *Base case* ($h = 0$): $m(0) = 1$ and $F(2) = F(0) + F(1) = 1 + 1 = 2$
— so the base case holds 😊
- *Inductive step* ($h > 0$): $m(h) = 1 + m(h-1) + m(h-2)$ Fails for $h = 1$
 $= 1 + F(h+1) - 1 + F(h) - 1$ by the inductive hypothesis
 $= F(h+2) - 1$ by the definition of F
😊

Fibonacci numbers and the Golden Ratio

- It turns out that

$$F(h) = \frac{\varphi^{h+1} - (1 - \varphi)^{h+1}}{\sqrt{5}} \quad \text{where } \varphi = \frac{1+\sqrt{5}}{2} = 1.618\dots$$

is the golden ratio

- Thus for large h , $m(h) = F(h + 2) - 1 \geq \frac{1}{\sqrt{5}} 1.618^{h+3}$

- Thus, for an AVL tree with n nodes and height h

$$n \geq \frac{1}{\sqrt{5}} 1.618^{h+3} \text{ and hence}$$

$$\begin{aligned} h &\leq \log_{1.618}(\sqrt{5}n) - 3 \\ &< \log_{1.618} n \approx 1.44 \log_2 n \end{aligned}$$

- Thus, the height of an AVL tree can be about 44% more than optimal
 - find, insert and delete in $O(\log n)$ worst-case time

Dictionary ADT

- A **dictionary** stores (key, value) pairs where keys are **totally ordered**
 - e.g., integers, reals, vectors, strings, ... `int cmp(key_t, key_t)`
 - For a key k , a dictionary has at most one pair (k, v)

```
#ifndef DICT_H
#define DICT_H
typedef void* DictEntry;
void setValue(DictEntry, value_t);
value_t getValue(DictEntry);
```

```
typedef void* Dict;
void freeDict(Dict);
int isEmpty(Dict);
int size(Dict);
Dict insert(Dict, key_t, value_t);
DictEntry get(Dict, key_t);
Dict delete(Dict, key_t);
DictEntry forall(Dict, int(*task)(int, void*), void*);
#endif
```

Assuming `cmp` runs
in $O(1)$ time!

Linked list	AVL tree
$O(1)$	$O(1)$
$O(n)$	$O(n)$
$O(n)$	$O(\log n)$

Replace old **value** if
key already exists


Research Note

- A comparison of dictionary implementations by Mark Neyer, 2009
- This paper explores several implementations of dictionaries:
 - hash tables, Red-Black Trees, AVL Trees, and Skip Lists
- The Microsoft .NET platform was chosen for this experiment because it is of interest to the author's experience in industry, and due to the availability of a precise memory measurement debugging library, namely `sos.dll`
- Implementations
 - `System.Collections hashtable`
 - AVL tree implementation: www.vcskicks.com
 - Red-Black tree implementation: www.devx.com/DevX/Article/36196
 - Skip List implementation: www.codeproject.com/KB/recipes/skiplist1.aspx

No mention of hardware!


Results (1/3)

- Size (in bytes) vs. n




Results (2/3)

- Creation time (in milliseconds) vs. n



Results (3/3)

- Query time for $100n$ keys (in milliseconds) vs. n



Priority Queues

- A **priority queue** is an ADT for storing (priority, data) pairs, where priorities are totally ordered (e.g., integers, strings, etc.)
 - Two variants (silly English):
 - low-magnitude priority means “high priority”
 - high-magnitude priority means “high priority”

```
#ifndef PQUEUE_H
#define PQUEUE_H
typedef void* PQEntry;
typedef void* PQueue;
int size(PQueue);
PQEntry insert(PQueue *pq, int priority, int data);
int deleteNext(PQueue pq, int *priority, int *data);
PQEntry update(PQueue pq, PQEntry e, int newPriority);
#endif
```


Priority Queue Implementations


	Linked List	Ordered Linked List	Ordered Binary Tree	Balanced Ordered Binary Tree
size	$O(1)$	$O(1)$	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$	$O(n)$	$O(\log n)$
deleteNext	$O(n)$	$O(1)$	$O(n)$	$O(\log n)$
update	$O(1)$	$O(n)$	$O(n)$	$O(\log n)$

- **Interview question:** All n (priority, data) items known in advance, only need deleteNext. What is the best implementation?
- AVL trees: pointer-based, large constant hidden in the $O(\cdot)$
- We will develop a fast, array-based implementation
 - size in $O(1)$ time
 - insert and deleteNext in $O(\log n)$ time [smaller constant]

Complete Binary Trees

- A binary tree is **complete** if all levels (except possibly the last) are completely filled, and all nodes are as far left as possible

▪ *Examples:*



- **Fact:** A complete binary tree with n nodes has height exactly $\lfloor \log_2 n \rfloor$
- Array representation:
 - root at index 0
 - for any node at index i :
 - left child at index $2i + 1$
 - right child at index $2i + 2$
 - parent at index $\lfloor (i - 1) / 2 \rfloor$ (unless $i = 0$)

Heaps

- **Definition:**

- A **minheap** is a complete binary tree in which every node's priority is less than or equal to the priority of its children (min. at root)
- A **maxheap** is a complete binary tree in which every node's priority is greater than or equal to the priority of its children (max. at root)

- We implement a heap as an array

- **Question:** Which of these arrays represent minheaps?

A

1	2	3	4	5	6
---	---	---	---	---	---

B

1	2	4	5	6	3
---	---	---	---	---	---

C

1	3	4	2	6	5
---	---	---	---	---	---

D

1	4	2	6	5	3
---	---	---	---	---	---

Sift Up

- Suppose the heap is implemented as a **struct** with array **a[]** where each array is a **struct** with priority **p** and data **d**

```
void siftUp(Heap *H, int i) { // minheap
    while(i > 0) {
        int j = (i-1)/2; // parent
        if(H->a[i].p >= H->a[j].p) { return; }
        SWAP(H->a, i, j);           Why not: SWAP(H->a[i], H->a[j])
        i = j;                     Temporary?
    }
}
```

- What does `siftUp(H, 3)` do if `H->a` =

1	3	4	2	6	5
---	---	---	---	---	---
- Insert = append to array + sift up the last element: $O(\log n)$ time

Sift Up

- Suppose the heap is implemented as a **struct** with array **a[]** where each array is a **struct** with priority **p** and data **d**

```
void siftUp(Heap *H, int i) { // minheap
    while(i > 0) {
        int j = (i-1)/2; // parent
        if(H->a[i].p >= H->a[j].p) { return; }
        SWAP(H->a, i, j);
        i = j;
    }
}
```

- Insert = append to array + sift up the last element: $O(\log n)$ time
- Time to “heapify” n items: $O(n \log n)$

Sift Down

```
void siftDown(Heap *H, int i) { // minheap
    while(i < (H->size)/2) { // i >= size/2 is leaf!
        int j = 2*i+1; // left child
        if(j < H->size-1 && H->a[j].p > H->a[j+1].p)
            j++; // right child
        if(H->a[i].p <= H->a[j].p) { return; }
        SWAP(H->a, i, j);
        i = j;
    }
}
```

- What does `siftDown(H, 0)` do if `H->a` =

7	2	4	3	6	5
---	---	---	---	---	---
- `deleteNext` = swap root with last, `size--`, `siftDown(0)`: $O(\log n)$ time
– update?

Heapify in $O(n)$ time

```
void heapify(Heap *H) { // minheap
    for(int i = H->size/2-1; i >= 0; i--) {
        siftDown(H, i);
    }
}
```

Sloppy

- Runtime analysis: $O(n)$ iterations, each takes $O(\log n)$ time: $O(n \log n)$
- Careful runtime analysis?
- Heapsort
- Two advanced data structures: Treaps and Fibonacci heaps

B-trees

Organization and main

R Bayer, E McCreight - Software

Summary Organization and main
considered. It is assumed that
backup store like a disc or a d

Cited by 2



References:

- Adelson-Velskii, G. M. and Landis, E. M. An Information Organization Algorithm. DANSSSR, No. 2, 1962.
- Foster, C. C. Information Storage and Retrieval Using AVL Trees. *Proc. ACM 20th Nat'l. Conf.* (1965), pp. 192-205.
- Gladun, V. P. Storage Organization for Key Search and Recording. *Cybernetics*, Vol. 1, No. 4, August 1965.
- Landauer, W. I. The Balanced Tree and Its Utilization in Information Retrieval. *IEEE Trans. on Electronic Computers*, Vol. EC-12, No. 6, December 1963.
- Sussenguth, E. H., Jr. The Use of Tree Structures for Processing Files. *Comm. ACM*, Vol. 6, No. 5, May 1963.

■ Actually

Organizat
access file is
on some pseudo
index organiza
of keys in time
index and k
formance of th
least 50% but
in a special d
performance bo

Large amount of data

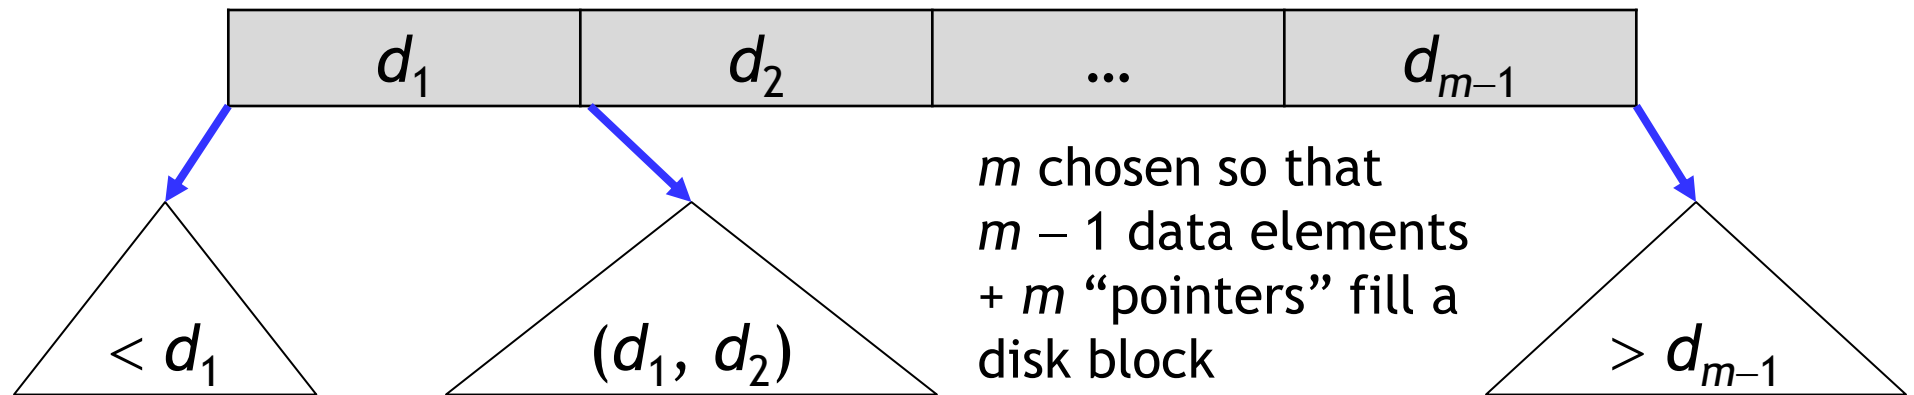
- Need to store most of the data on disk (or in the cloud)
 - getting data from these sources is MUCH slower than RAM

Data from 2014	Component	Latency
	1 CPU cycle	0.3 ns
	L1 cache access	0.9 ns
	L2 cache access	2.8 ns
	L3 cache access	12.9 ns
	RAM access	120 ns
	SSD access	50-150 μ s
	HDD access	1-10 ms
	SF to NYC internet	40 ms

- **Key idea:** If we have to go to the disk, make it count!
 - each disk access fetches a whole **block** of data (a few KB)

Nodes in a B-tree


- In a B-tree of order m , each node has up to m children



- The root is either a leaf, or has at least two children
- Each internal node except the root has between $\lceil m/2 \rceil$ and m children
- All leaves are at the same level
 - tree height is $\approx \log_m n$
 - m is of the order of 10^2 or more, so for $n \approx 10^9$, depth is around 5

Operations on B-trees

- Find is a generalization of Find for ordered binary trees
 - binary search used to find the correct pointer to follow



- To insert d , find the leaf node where it should have been
 - insert into that node (in sorted order)
 - if the node was full (i.e., it had $m - 1$ elements), promote the middle element to the parent-level and split this into two half-filled nodes
 - promote/split may cascade up to the root (tree height will grow)
- Delete is similar: remove data, then try to compress with adjacent sibling

An optimization: B⁺ trees

- Suppose data is made up of (small) keys and (large) values
 - search done on the basis of keys
- Internal nodes only store keys
 - values are stored only at leaves
- Thus, a single disk block can accommodate many more keys
 - bushier tree \Rightarrow shallower tree
- Insert animation from OpenDSA

B/B⁺ trees: Frequent nodes

- In both B-trees and B⁺-trees, the root is processed frequently
- While the data structure is needed, the root can be stored in RAM
 - the root can be written out to the disk when the data structure is no longer needed by the program (e.g., when the program terminates)
- RAM is large enough to accommodate more than just the root
 - but it cannot accommodate everything
- *Note:* This is exactly the problem of **caching**
 - cache faster than RAM (just like RAM faster than disk)
 - cache smaller than RAM (just like RAM smaller than disk)
- **Key Question:** What to save and what to **evict**?

Eviction Strategies

- FIFO: First In First Out
 - data structure?
- LRU: Least Recently Used
 - data structure?
- FF: Furthest in Future
 - provably optimal
 - data structure?
- Aside: Splay trees
 - whenever you insert or find something, **splay** it to the top
 - $O(\log n)$ amortized find/insert/delete

Questions

1. In a **full** tree, every node has either 0 or 2 children
 - What is the maximum height of a full tree with n nodes
2. Given pointers p and q to two nodes in an balanced **binary search tree** (ordered binary tree) with n nodes, describe a fast algorithm to find the lowest common ancestor of p and q
 - Answer the same question, assuming that the tree is unordered and not necessarily balanced
3. In a **general tree**, nodes can have an arbitrary number of children
 - Describe a space-efficient way to implement general trees in C